

# Lab 5:

## OpenMP Part 2

Jan Izquierdo  
jan.izquierdo@

Laura LLorente  
laura.llorente@

17/05/2024

**Load the necessary modules:**

```
module load openmpi/4.1.1
module load gcc/12.1.1
```

**1 node running a single MPI process with 12 threads.**

```
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --ntasks-per-node=1
----module load gcc/12.1.1
export OMP_NUM_THREADS= 12
mpirun hellohybrid
```

**2 nodes running a single process in each node and each process is divided into 12 threads**

```
#SBATCH --nodes=2
#SBATCH --ntasks=2
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=12
----
export OMP_NUM_THREADS=12
mpirun hellohybrid
```

**2 nodes running two processes each, with 6 threads per process**

```
#SBATCH --nodes=2
#SBATCH --ntasks=4
#SBATCH --ntasks-per-node=2
#SBATCH --cpus-per-task=6
----
export OMP_NUM_THREADS=6
mpirun hellohybrid
```

**2 nodes running 24 single-threaded MPI processes (12 per node)**

```
#SBATCH --nodes=2
#SBATCH --ntasks=24
#SBATCH --ntasks-per-node=12
----
export OMP_NUM_THREADS=12
mpirun hellohybrid
```

## 1. Modify the program to count ALL 0-sum subarrays that exist in a given list of integers

To count the number of sums equal to 0 we modified the existing function to add 1 to a counter instead of finishing with a return every time a sum equal to 0 was found:

```
int sumCounter(int arr[], int n){
    int c=0;

    for (int i=0;i<n;i++){

        int sum = arr[i]; //if element is 0 -> +1 to counter
        if (sum == 0){c++;}

        for (int j = i + 1; j < n; j++) { //starting from current position and forward
            sum += arr[j]; //every time you advance a position add the corresponding number
            if (sum == 0) // if the sum of the numbers from i so far is 0, add 1 to the counter
                c++;
        }
    }
    return c;
}
```

## 2. Parallelize the function that counts the number of 0-sum subarrays using the OpenMP library. The program must have two versions of the same function: the sequential and the parallel, to verify that the program provides the same correct results in both cases

In order to parallelize the sumCounter, we added the directive **#pragma omp parallel for reduction(+:c)** in order to execute each for loop with a different thread, the reduction clause gives each thread a private copy of the c counter and at the end of the parallel region, all the local copies are combined into the final result :

```
int PsumCounter(int arr[], int n){ //same reasonings as sumCounter
```

```

int i, j, c=0; //i, j defined here for parallels later

#pragma omp parallel for reduction(+:c) //distributes an i for
each process, makes c private and adds them up after the for loop
is finished to ensure its correct
for (i=0;i<n;i++){
    int sum = arr[i];
    if (sum == 0){c++;}

    #pragma omp parallel for reduction(+:c) // distributes a j
for each process, makes c private and adds them up when the loop
is finished to ensure its correct
    for (j = i + 1; j < n; j++){
        sum+=arr[j];
        if (sum==0){c++;}
    }
}
return c;
}

```

In order to test both functions, we used different arrays to check if the results were right. In the final test we set different seeds with a large array (2000). The result proves that both sequential and parallel functions always prints the same correct result:

```

[biohpc-25@clus-login LAB5]$ ./sum0

***** Found a subarray with 0 sum *****

**** Found 8650 subarrays with 0 sum using a sequential program ****

***** Found 8650 subarrays with 0 sum using a parallel program *****

```